

SIEIS — Foolproof Implementation Plan

From Current State → Full ML-Powered Dashboard

Date: February 19, 2026 | For: Aishwarya (Group 8, AAI-530)

Where You Are Right Now (The Foundation)

Think of your project like building a house. You've already laid the **foundation and plumbing**:

DONE (Your Foundation)

- Docker Compose stack running (InfluxDB, MinIO, Kafka/Redpanda, Simulator, Consumer)
- Historical data loader (1.3M+ records → InfluxDB + MinIO)
- Real-time simulator (42 motes emitting sensor data through Kafka)
- Consumer writing to both InfluxDB (hot) and MinIO (cold) storage
- Verification scripts for both data stores
- Project structure, .env config, requirements.txt, tests folder
- Documentation (architecture, roadmap, deployment, ML docs)

TO BUILD (This Plan)

- Phase A: Data Validation (the "health check")
- Phase B: FastAPI Backend (the "brain")
- Phase C: ML Model Training (the "intelligence")
- Phase D: ML Inference API (the "prediction engine")
- Phase E: Streamlit Dashboards (the "face")
- Phase F: Incremental Retraining (the "learning loop")
- Phase G: Integration & Polish (the "finishing touches")

Beginner Concepts You Need to Know

What is an ML Pipeline?

Analogy: A Factory Assembly Line

Imagine a chocolate factory:

1. **Raw cocoa beans arrive** → That's your raw sensor data in MinIO/InfluxDB
2. **Beans are cleaned, roasted, sorted** → That's *feature engineering* (preparing data for the model)
3. **A machine molds chocolates** → That's *model training* (the algorithm learns patterns)
4. **Quality inspector checks each chocolate** → That's *model evaluation* (checking accuracy)
5. **Chocolates are boxed and shipped** → That's *model deployment* (serving predictions via API)

6. **Customer feedback improves the recipe** → That's *incremental retraining* (model gets smarter over time)

What is Isolation Forest?

Analogy: The Odd-One-Out Game

Imagine you have 100 apples. 95 are red, 5 are blue. If you randomly pick features to split them, the blue ones get "isolated" very quickly (in just 1-2 splits), while red ones take many splits to separate. Isolation Forest works the same way — anomalies are data points that are easy to isolate. Short path in the tree = anomaly. Long path = normal.

Your sensor data has temperature, humidity, light, and voltage. If a mote suddenly reads 50°C when everything else is 22°C, Isolation Forest catches it because that reading gets isolated in very few splits.

What is Incremental Retraining?

Analogy: A Student Studying Weekly

Instead of cramming everything the night before an exam (training once on ALL data), the student studies a little each week (retraining on new batches). The model doesn't forget old knowledge — it gets updated with fresh patterns from recent data. In your case, you retrain manually whenever you load new sensor data.

Hot Path vs. Cold Path

Analogy: Your Kitchen vs. Your Pantry

- **Hot Path (InfluxDB):** Like your kitchen counter — recent items you use constantly. Fast access, limited space. Holds the last 30 days of sensor data for real-time monitoring.
 - **Cold Path (MinIO):** Like your pantry — everything stored long-term in organized containers (Parquet files). Holds ALL historical data for training ML models and trend analysis.
-

Phase A: Data Validation (2–3 hours)

What: Make sure the data you loaded is clean and complete before building anything on top. **Why:** Building an ML model on bad data is like baking a cake with spoiled ingredients — the output will be garbage.

What to Build

One script: `scripts/validate_data_quality.py`

This script connects to BOTH InfluxDB and MinIO and checks:

Check	What It Means	Pass Criteria
Total record count	Did everything load?	1,500,000+ records
Duplicate ratio	Are there repeated readings?	Less than 1% duplicates
Mote coverage	Are all 42 sensors present?	42 unique mote IDs
Date range	Does the data span the right period?	Jan 19 – Feb 19, 2026
Value ranges	Are readings physically plausible?	Temp: -10 to 50°C, Humidity: 0-100%, Light: 0-500, Voltage: 2.0-3.0V
Missing values	How many gaps exist?	Less than 5% missing
InfluxDB vs MinIO match	Do both stores agree?	Record counts within 1% of each other

How to Build It in VS Code

1. Create the file at `scripts/validate_data_quality.py`
2. It should import `influxdb_client` and `minio` from your existing config
3. Run each check, collect pass/fail results
4. Print a summary report at the end

Expected Output

DATA QUALITY VALIDATION REPORT

✓ Total Records: 1,500,000
 ✓ Duplicates: 2,340 (0.16%) — PASS
 ✓ Unique Motes: 42/42 — PASS
 ✓ Date Range: 2026-01-19 to 2026-02-19 — PASS
 ✓ Missing Values: 45,000 (3%) — PASS
 ✓ InfluxDB ↔ MinIO Consistency — PASS

Final: ✓ PASSED (99.8% quality)

Done When

- Script runs without errors
- All checks pass
- You can confidently say "my data is clean"

Phase B: FastAPI Backend (8–10 hours)

What: Build a REST API so your dashboards and ML models can talk to your data stores. **Why:** Think of FastAPI as the "waiter" in a restaurant. The dashboard (customer) asks for data, the API (waiter) goes to the kitchen (InfluxDB/MinIO), gets what's needed, and serves it back neatly.

Folder Structure to Create

```
src/app/api/
├── __init__.py      ← Makes this a Python package (empty file)
├── main.py          ← The "restaurant entrance" — starts the FastAPI app
└── routes/
    ├── __init__.py
    ├── sensor_data.py   ← Endpoints for sensor readings (GET latest, GET history)
    ├── analytics.py     ← Endpoints for daily summaries, aggregates
    └── ml.py            ← Endpoints for ML predictions (built in Phase D)
    ├── schemas.py       ← Defines the shape of requests/responses (like a menu)
    ├── dependencies.py  ← Shared connections to InfluxDB and MinIO
    └── utils.py         ← Helper functions (date parsing, formatting)
```

Key Endpoints to Build (in order)

Round 1 — Health & Basics (build these first):

Endpoint	Purpose	Beginner Note
GET /api/v1/health	Check if API + databases are up	Always build this first — it's your "is everything alive?" check
GET /api/v1/sensors/latest	Latest reading from all 42 motes	Returns one row per mote with the most recent values
GET /api/v1/sensors/{mote_id}/latest	Latest reading for one mote	Same as above but filtered to one sensor

Round 2 — Time-Series Queries:

Endpoint	Purpose	Beginner Note
GET /api/v1/sensors/readings?mote_id=1&start=...&end=...	Historical readings for a mote	This powers your trend charts in the dashboard
GET /api/v1/sensors/{mote_id}/stats?start=...&end=...	Aggregated stats (avg, min, max, stddev)	Used for summary cards and KPI tiles

Round 3 — Analytics:

Endpoint	Purpose
GET /api/v1/analytics/daily-summary?date=2026-02-19	One-day overview across all motes
GET /api/v1/motes	List all 42 motes with status

How to Build in VS Code

1. Install dependencies: `pip install fastapi uvicorn pydantic`
2. In `main.py`: Create the FastAPI app, add CORS middleware, include route files
3. In `dependencies.py`: Create InfluxDB and MinIO client instances using `.env` variables
4. In `schemas.py`: Define Pydantic models for your request/response shapes
5. Build routes one at a time, test each with the auto-generated Swagger docs at <http://localhost:8000/docs>
6. Run with: `uvicorn src.app.api.main:app --reload`

How to Test Each Endpoint

FastAPI auto-generates interactive documentation. Open <http://localhost:8000/docs> in your browser, click any endpoint, hit "Try it out", and see the response live. No Postman needed.

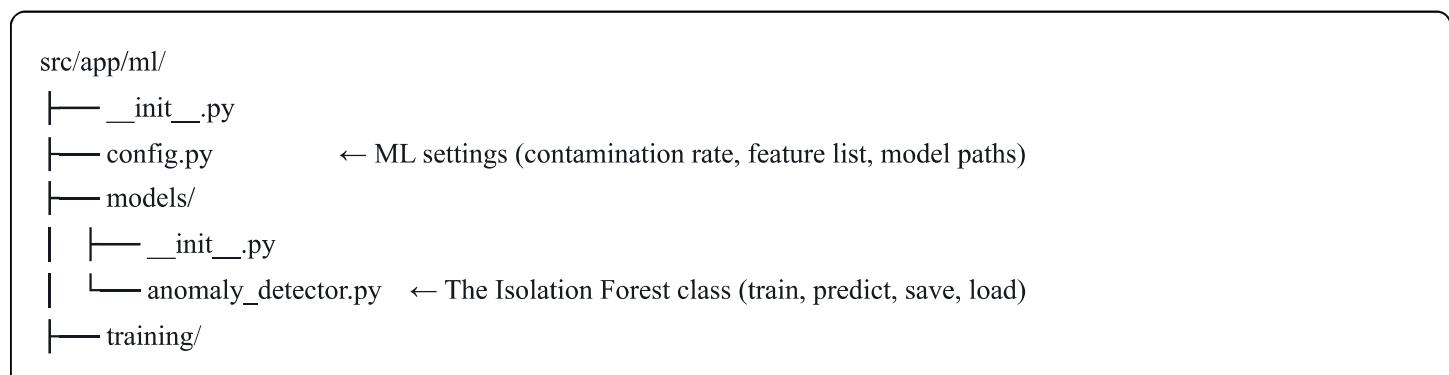
Done When

- GET /health returns `{"status": "healthy"}`
- GET /sensors/latest returns readings for all 42 motes
- GET /sensors/readings returns historical data for any mote + date range
- Swagger docs render at `/docs`

Phase C: ML Model Training (6–8 hours)

What: Train an Isolation Forest anomaly detection model on your historical sensor data. **Why:** This is the "intelligence" of your system — it learns what "normal" looks like so it can flag unusual readings.

Folder Structure to Create



```

|   └── __init__.py
|   └── data_loader.py      ← Pull data from MinIO Parquet files into a pandas DataFrame
|   └── feature_engineer.py  ← Create features (hour_of_day, day_of_week, rolling averages)
└── serving/
    └── __init__.py
        └── predictor.py     ← Load a saved model and make predictions

models/           ← Saved model files (at project root)
    └── anomaly_detector_v1_20260219.pkl
    └── model_registry.json  ← Tracks which model version is "current"

scripts/
    └── train_model.py     ← One-click training script

```

The Training Pipeline (Step by Step)

Think of this as a **cooking recipe with 5 steps**:

Step 1 — Gather Ingredients (Data Loading)

- Source: MinIO Parquet files (cold storage, full historical coverage)
- Read all Parquet files into one big pandas DataFrame
- Expect ~1.3M rows with columns: `timestamp, mote_id, temperature, humidity, light, voltage`

Step 2 — Prep the Ingredients (Feature Engineering)

- From the timestamp, extract: `hour_of_day` (0-23) and `day_of_week` (0-6)
- These matter because sensor patterns change by time of day (lights off at night, temperature drops)
- Optional but valuable: rolling mean and rolling standard deviation over the last 6 readings per mote
- Final feature set: `temperature, humidity, light, voltage, hour, day_of_week` (6 features)
- Handle missing values: drop rows where any sensor reading is null

Step 3 — Cook (Model Training)

- Algorithm: `IsolationForest` from scikit-learn
- Key parameters:
 - `contamination=0.05` → Tells the model "expect about 5% of readings to be anomalies"
 - `n_estimators=100` → Number of isolation trees (more = better but slower; 100 is a good default)
 - `random_state=42` → Makes results reproducible (same data → same model every time)
- Before training: scale features using `StandardScaler` (converts all values to similar ranges so no single feature dominates)
- Fit the model on the scaled historical data

Step 4 — Taste Test (Model Evaluation)

- Run predictions on the training data itself
- Check: What percentage did it flag as anomalies? (Should be close to 5%)
- Log training metrics: number of samples, features, anomaly ratio, training time
- If anomaly ratio is wildly off (like 30%), your contamination parameter needs adjusting

Step 5 — Package and Label (Save Model)

- Save the trained model + scaler using `[joblib]` (a Python library for saving ML objects)
- Filename format: `[models/anomaly_detector_v1_YYYYMMDD.pkl]`
- Update `[models/model_registry.json]` with the new version, timestamp, and metrics
- This registry tells the API which model file to load

What the Training Script Does

`[scripts/train_model.py]` orchestrates all 5 steps:

1. Load data from MinIO
2. Engineer features
3. Train Isolation Forest
4. Evaluate and print metrics
5. Save model artifact + update registry

Run it with: `python scripts/train_model.py`

Expected Training Output

SIEIS ML Training Pipeline

```
Loading data from MinIO... 1,312,456 records loaded
Engineering features...     6 features created
Handling missing values... 1,280,000 clean records
Scaling features...        StandardScaler fitted
Training Isolation Forest... Done in 12.3 seconds
Evaluation:
Samples:    1,280,000
Features:   6
Anomalies:  64,000 (5.0%)
Model size: 8.2 MB
Saving model...      models/anomaly_detector_v1_20260219.pkl
Registry updated.
```

Training complete!

Done When

- `scripts/train_model.py` runs end-to-end without errors
- A `.pkl` file exists in `models/`
- `model_registry.json` has the correct version info
- Anomaly ratio is between 3-8% (reasonable for sensor data)

Phase D: ML Inference API (4–5 hours)

What: Add prediction endpoints to your FastAPI app so dashboards can send sensor readings and get back "normal" or "anomaly". **Why:** The trained model sitting in a file is useless until something can ask it questions. The API makes the model accessible.

Analogy: Training the model is like teaching a dog to detect drugs. Deploying it is like bringing the dog to the airport and letting passengers walk past it.

What to Build

1. Predictor Service (`(src/app/ml/serving/predictor.py)`)

- On startup: reads `model_registry.json`, loads the latest `.pkl` file
- Exposes a `predict()` method that takes sensor values and returns anomaly score + label
- Keeps the model in memory (no file I/O per prediction = fast)

2. ML API Route (`((src/app/api/routes/ml.py))`)

Endpoint	Method	Purpose
<code>/api/v1/ml/predict/anomaly</code>	POST	Send one reading → get anomaly score
<code>/api/v1/ml/predict/batch</code>	POST	Send multiple readings → get scores for all
<code>/api/v1/ml/model/info</code>	GET	Returns current model version, training date, metrics
<code>/api/v1/ml/retrain</code>	POST	Triggers a retrain (Phase F)

How a Prediction Works (End to End)

Dashboard sends:

POST /api/v1/ml/predict/anomaly

Body: {"temperature": 45.2, "humidity": 30.1, "light": 250, "voltage": 2.8}

API receives → extracts features → adds hour/day_of_week → scales with saved scaler
→ runs through Isolation Forest → gets anomaly score

API returns:

```
{  
    "is_anomaly": true,  
    "anomaly_score": -0.72,  
    "severity": "high",  
    "details": "Temperature 45.2°C significantly above normal range (18-28°C)"  
}
```

How Severity Works

Score Range	Label	Meaning
-0.3 to 0.0	Low	Slightly unusual but likely fine
-0.6 to -0.3	Medium	Worth investigating
Below -0.6	High	Definitely anomalous

(These thresholds are tunable in `src/app/ml/config.py`)

How to Build in VS Code

1. Create `predictor.py` — a class that loads model on init, has a `predict()` method
2. In `dependencies.py` — instantiate the predictor as a singleton (loaded once, shared across requests)
3. Create `routes/ml.py` — define the endpoints, call predictor methods
4. Register the ML routes in `main.py`
5. Test via Swagger: POST a sample reading, verify you get an anomaly response

Done When

- `POST /ml/predict/anomaly` returns a valid prediction
- `GET /ml/model/info` shows model version and training metrics
- Prediction latency is under 50ms per reading
- Batch endpoint handles 100+ readings at once

Phase E: Streamlit Dashboards (10–12 hours)

What: Build two visual dashboards — one for real-time monitoring (InfluxDB) and one for historical analysis (MinIO). **Why:** Dashboards are how humans consume data. All the infrastructure you've built becomes useful when someone can actually see trends, anomalies, and patterns on a screen.

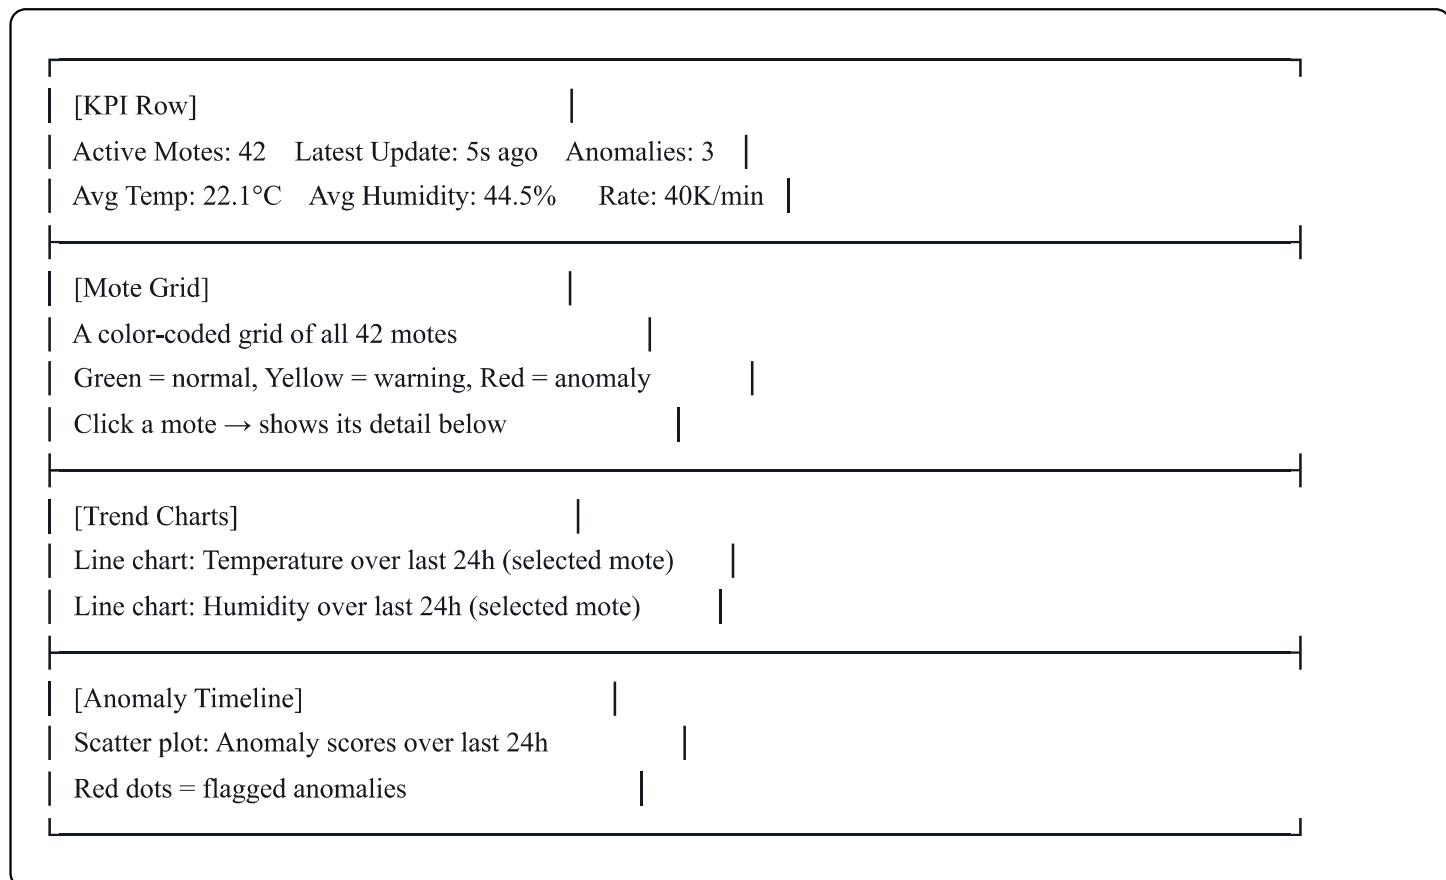
Folder Structure to Create

```
src/app/dashboard/
    ├── app.py           ← Main Streamlit entry point with sidebar navigation
    └── pages/
        ├── 1_realtime.py   ← Hot path dashboard (InfluxDB data)
        ├── 2_historical.py  ← Cold path dashboard (MinIO data)
        ├── 3_ml_insights.py  ← ML predictions and anomaly analysis
        └── 4_system_health.py  ← Pipeline monitoring (Kafka lag, DB status)
    └── utils/
        ├── __init__.py
        ├── influx_client.py  ← Helper to query InfluxDB for dashboard
        ├── minio_client.py  ← Helper to read MinIO Parquet for dashboard
        └── api_client.py  ← Helper to call your FastAPI endpoints
```

Dashboard 1: Real-Time Monitoring (pages/1_realtime.py)

Data Source: InfluxDB (last 24 hours) **Refresh:** Auto-refresh every 10-30 seconds using Streamlit's `st.rerun()`

Layout:



Dashboard 2: Historical Analysis (pages/2_historical.py)

Data Source: MinIO Parquet files (full date range) **Refresh:** On-demand (user selects date range)

Layout:

[Filters]

Date Range Picker: [Jan 19] to [Feb 19] Mote: [All ▼]

[KPI Row]

Total Records: 1.3M Days Covered: 31 Avg Quality: 99%

[Daily Trends]

Line chart: Daily average temperature across all motes

Line chart: Daily average humidity

[Seasonality Heatmap]

X-axis: hour of day (0-23)

Y-axis: day of week (Mon-Sun)

Color: average temperature → reveals daily/weekly patterns

[Distribution Histograms]

4 histograms: temperature, humidity, light, voltage

Shows the spread of values across the entire history

Dashboard 3: ML Insights (pages/3_ml_insights.py)

Data Source: Your FastAPI ML endpoints **Purpose:** Show what the model is finding

Layout:

[Model Info Bar]

Model: Isolation Forest v1 Trained: Feb 19 Status:

[Live Anomaly Feed]

Table: Recent anomalies (timestamp, mote, metric, severity)

[Anomaly Score Distribution]

Histogram of scores: most readings cluster near 0 (normal)

Tail stretching to -1 (anomalous)

[Manual Prediction Tool]

4 sliders: temperature, humidity, light, voltage

"Predict" button → calls API → shows result

[Retrain Button]

"Retrain Model" → triggers /api/v1/ml/retrain

Shows progress and new model metrics when done

How to Build in VS Code

1. Install: `pip install streamlit plotly`
2. Start with `app.py` — just the sidebar navigation and page routing
3. Build `1_realtime.py` first (most visible, most impressive for a demo)
4. Use `st.metric()` for KPI cards, `plotly.graph_objects` for charts
5. Query InfluxDB directly from the dashboard OR call your FastAPI endpoints via `requests`
6. Run with: `streamlit run src/app/dashboard/app.py`
7. Opens at `http://localhost:8501`

Done When

- All 4 pages load without errors
 - Real-time page auto-refreshes and shows live data
 - Historical page renders trend charts for any selected date range
 - ML page shows anomaly predictions and model info
 - Charts render in under 2 seconds
-

Phase F: Incremental Retraining (3–4 hours)

What: A script + API endpoint that retrains the model on new data without starting from scratch. **Why:** Your simulator keeps generating new data daily. The model trained on Jan 19 – Feb 19 data won't know about Feb 20+ patterns. Retraining incorporates new patterns.

Analogy: Updating a Textbook Edition You don't rewrite the whole textbook — you add a new chapter with recent discoveries and update the edition number. Similarly, incremental retraining pulls recent data, re-fits the model, saves a new version, and swaps it into the running API.

What to Build

1. Retrain Script (`scripts/retrain_model.py`)

This script does:

1. Read the current model version from `model_registry.json`
2. Pull the last 7–30 days of data from MinIO (configurable)
3. Engineer features (same pipeline as initial training)
4. Train a NEW Isolation Forest on this data
5. Save as `models/anomaly_detector_v{N+1}_YYYYMMDD.pkl`
6. Update `model_registry.json` to point to the new version
7. Print comparison: old model metrics vs new model metrics

2. API Endpoint (`POST /api/v1/ml/retrain`)

When called:

- Runs the retrain pipeline in the background
- Returns immediately with `{"status": "retraining_started", "estimated_time": "30s"}`
- On completion: the predictor service automatically reloads the new model
- Next prediction request uses the updated model

3. Model Hot-Swap in Predictor

The predictor service checks `model_registry.json` periodically (every 60 seconds) or reloads on demand after retraining. This way you don't need to restart the API server.

When to Retrain

Trigger	How
After loading new data	Run <code>python scripts/retrain_model.py</code> manually
Weekly maintenance	Run the script every Monday
From dashboard	Click "Retrain Model" button on ML Insights page
Via API	<code>POST /api/v1/ml/retrain</code>

Done When

- `scripts/retrain_model.py` creates a new model version
- `model_registry.json` updates correctly
- API's next prediction uses the new model without restart
- Dashboard "Retrain" button works end-to-end

Phase G: Integration & Polish (4–5 hours)

What: Wire everything together, add tests, and make it demo-ready.

Integration Checklist

Task	What It Means
End-to-end data flow test	Simulator → Kafka → Consumer → InfluxDB + MinIO → API → Dashboard (all connected)
ML pipeline test	MinIO data → Training → Model saved → API loads it → Dashboard shows predictions

Task	What It Means
Error handling	What happens when InfluxDB is down? When model file is missing? Add graceful fallbacks
Logging	Add <code>print()</code> or Python <code>logging</code> statements at key points so you can debug issues
Requirements update	Make sure <code>requirements.txt</code> has ALL packages: <code>fastapi</code> , <code>unicorn</code> , <code>streamlit</code> , <code>plotly</code> , <code>scikit-learn</code> , <code>joblib</code> , <code>pandas</code> , <code>pyarrow</code> , <code>influxdb-client</code> , <code>minio</code> , <code>pydantic</code> , <code>python-dotenv</code>

Test Scenarios to Verify

1. **Start fresh:** `(docker-compose down && docker-compose up -d)` → wait 30s → run `(verify_influxDb.py)` → data appears
2. **Train model:** `(python scripts/train_model.py)` → `(.pkl)` file created → metrics printed
3. **Start API:** `(unicorn src/app/api/main:app --reload)` → `(/docs)` loads → `(/health)` returns OK → `(/ml/model/info)` shows model version
4. **Test prediction:** POST to `(/ml/predict/anomaly)` with sample values → get anomaly score back
5. **Start dashboard:** `(streamlit run src/app/dashboard/app.py)` → all 4 pages render → charts show data
6. **Retrain:** Click "Retrain" on dashboard → new model version appears → predictions use new model

Final Project Structure

```

SIEIS/
├── .env           ← Environment variables (tokens, URLs)
├── docker-compose.yml   ← Infrastructure (InfluxDB, MinIO, Kafka, Simulator, Consumer)
├── requirements.txt    ← All Python dependencies
├── Dockerfile.consumer    ← Consumer container
├── Dockerfile.simulator    ← Simulator container
├── data/           ← Raw datasets
├── models/          ← Saved ML models (.pkl files)
│   ├── anomaly_detector_v1_20260219.pkl
│   └── model_registry.json
└── scripts/
    ├── load_historical_data.py  ← Phase 1 (done)
    ├── verify_influxDb.py      ← Phase 1 (done)
    ├── verify_minio_storage.py  ← Phase 1 (done)
    ├── validate_data_quality.py ← Phase A
    ├── train_model.py         ← Phase C
    └── retrain_model.py       ← Phase F
└── src/app/
    ├── simulator/           ← Phase 1 (done)
    ├── consumer/            ← Phase 1 (done)
    ├── config/              ← Phase 1 (done)
    └── api/                 ← Phase B + D

```

```
    |   |   └── main.py
    |   └── routes/
    |       ├── sensor_data.py
    |       ├── analytics.py
    |       └── ml.py
    └── schemas.py
    └── dependencies.py
    └── utils.py
└── ml/           ← Phase C + D + F
    ├── config.py
    ├── models/
    |   └── anomaly_detector.py
    ├── training/
    |   ├── data_loader.py
    |   └── feature_engineer.py
    └── serving/
        └── predictor.py
└── dashboard/     ← Phase E
    ├── app.py
    ├── pages/
    |   ├── 1_realtime.py
    |   ├── 2_historical.py
    |   ├── 3_ml_insights.py
    |   └── 4_system_health.py
    └── utils/
        ├── influx_client.py
        ├── minio_client.py
        └── api_client.py
└── tests/
    ├── test_api_endpoints.py
    ├── test_ml_pipeline.py
    └── test_data_quality.py
└── Documentation/
    ├── ARCHITECTURE.md
    ├── ML_MODELS.md
    ├── ML_ACTION_PLAN.md
    ├── PROJECT_ROADMAP.md
    ├── DEPLOYMENT.md
    └── NEXT_STEPS.md
```

Master Timeline

Phase	Task	Hours	Depends On	Key Deliverable
A	Data Validation	2–3h	Data loaded <input checked="" type="checkbox"/>	Quality report passes
B	FastAPI Backend	8–10h	Phase A	Swagger docs + working endpoints
C	ML Model Training	6–8h	Phase A	Saved <code>.pkl</code> model file
D	ML Inference API	4–5h	Phase B + C	<code>/ml/predict/anomaly</code> works
E	Streamlit Dashboards	10–12h	Phase B + D	4 dashboard pages rendering
F	Incremental Retraining	3–4h	Phase C + D	Retrain button creates new model version
G	Integration & Polish	4–5h	All above	Full end-to-end demo works
TOTAL		37–47h		

Suggested Daily Schedule (If Working Full Days)

Day	Focus	Outcome
Day 1	Phase A + start Phase B	Data validated, API skeleton running
Day 2	Finish Phase B	All sensor/analytics endpoints working
Day 3	Phase C	Trained ML model saved
Day 4	Phase D + start Phase E	Prediction API working, real-time dashboard started
Day 5	Finish Phase E	All 4 dashboard pages rendering
Day 6	Phase F + Phase G	Retraining works, full integration test passes

VS Code Tips for This Project

1. **Open the project folder** as workspace root (the folder with `docker-compose.yml`)
2. **Use the integrated terminal** (`Ctrl+``) — run Docker, Python, and Streamlit from here
3. **Install Python extension** by Microsoft — gives you linting, autocomplete, debugging
4. **Run multiple terminals:** One for Docker logs, one for FastAPI, one for Streamlit
5. **Use `.env` file support:** Install "DotENV" extension for syntax highlighting

6. Debug FastAPI: Add a VS Code launch config for `uvicorn` with `--reload`

Terminal Layout While Developing

```
Terminal 1: docker-compose logs -f      ← Watch infrastructure  
Terminal 2: uvicorn src.app.api.main:app --reload  ← API server  
Terminal 3: streamlit run src/app/dashboard/app.py  ← Dashboard  
Terminal 4: (free for running scripts)  ← Training, validation, etc.
```

Common Pitfalls to Avoid

Pitfall	Why It Happens	How to Avoid
Model training fails with NaN errors	Missing values in data weren't handled	Always <code>dropna()</code> or <code>fillna()</code> before training
API returns 500 errors	Model file not found or wrong path	Use absolute paths and check <code>model_registry.json</code> on startup
Dashboard is slow	Querying all 1.3M records on every refresh	Use aggregation queries (hourly/daily averages) instead of raw data
Retraining overwrites old model	File naming doesn't include version	Always use timestamps in filenames: <code>model_v2_20260220.pkl</code>
Features mismatch between training and prediction	Training used 6 features but API sends 4	Save the feature list alongside the model and validate on prediction
Docker containers out of memory	InfluxDB + MinIO + Kafka eat RAM	Set memory limits in <code>docker-compose.yml</code> , restart containers if needed

This plan covers your complete path from current state to a fully working ML-powered IoT dashboard. Build each phase sequentially — each one builds on the previous. When in doubt, verify the previous phase is solid before moving forward.